



Vrije Universiteit Brussel

Faculty of Science
Department of Computer Science

Rate-monotonic scheduling

Operating systems & security

Filip Moons

Promotor: Prof. Dr. M. Timmerman

January 2014



Contents

1	Introduction	2
2	Real-time operating systems	3
3	Scheduling	4
3.1	Definitions	4
3.2	Scheduling algorithms	4
3.2.1	Static priority scheduling algorithms	4
3.2.2	Dynamic priority scheduling algorithms	4
3.2.3	Preemptive priority scheduling algorithms	5
3.2.4	Rate-monotonic scheduling	5
3.3	Schedulability tests	5
3.3.1	Liu & Layland lower bound	5
3.3.2	Response time (RT) test	7
3.4	Time-Demand function	7
3.5	Dynamic scheduling variant of RMS	9
4	Extensions	11
4.1	Non-periodic services	11
4.1.1	Sporadic service	11
4.2	Priority ceiling protocol	12
4.2.1	Priority inversion	12
4.2.2	The protocol	12
4.3	Tasks with $D_i \leq T_i$	14
4.4	Other extensions	14
5	Simulations using the schedulability analysis tool Cheddar	15
5.1	What is Cheddar?	15
5.2	Configuration	15
5.3	Simulation of RMS of example 5.3	15
6	Conclusion	16

1 Introduction

This paper gives a good overview of the basic concepts of Rate Monotonic scheduling. We first introduce real-time operating systems, because that's the field where RMS is applied. Second, we introduce the needed definitions to know before studying the algorithm. We make make a lot of assumptions about the tasks that must be executed, to become an easy configuration in which we can introduce RMS. Next, the algorithm is introduced. The second section closes with the schedulability tests used to decide if a task set can be scheduled by RMS or not. The next section is about extensions of the Rate Monotonic theory. It gives some examples of the theory applied on task sets with different assumptions than in section 1. These extensions are in fact generalizations of the theory because they often just abandon some of the assumptions. We close this paper with some concrete simulations of RMS in the program Cheddar.

2 Real-time operating systems

A real-time operating system (RTOS) is an operating systems that serves real-time application request. The operating system can process data as it comes in, normally without notable delay. Processing times are measured in milliseconds. A real-time operating system has an advanced algorithm for scheduling it's tasks. This paper is about such an advanced algorithm: rate-monotonic scheduling. The real-time operating systems using RMS are the oftene preemptive and have quite constant response times. Preemptive operating systems are capable of temporarily interrupting a task to give the priority to another task, without requiring cooperation of the task. The interrupted task is resumed at a later time.

3 Scheduling

3.1 Definitions

Definition 3.1. A task τ_i is a process or a thread that has to be periodically executed in a period T_i . The worst case execution time of a task τ_i is denoted as C_i , with $C_i \leq T_i$. The deadline D_i of τ_i is the available time on the processor to execute the task. In this paper however, we consider $D_i = T_i$. The priority of a task is indicated by its index number: τ_i has a higher priority than τ_j when $i < j$.

Definition 3.2. A job of task τ_i is a single period of task τ_i .

Definition 3.3. A task τ_i has a phasing I_i ($0 < I_i < T_i$). This phasing gives the time at which the first job of τ_i occurs.

Definition 3.4. A task set is a set of tasks, the indexes of the task indicate the priorities of these tasks relative to each other:

$$\mathcal{T} = \{\tau_1, \tau_2, \dots, \tau_n\}$$

Remark 3.5. In this section, we have to make some assumptions about tasks: switching between tasks is instantaneous, there is no interaction between tasks, jobs are ready to be executed at the beginning of their period and they quit the CPU after the job is finished. A task can only have one job that's being executed, new jobs have to wait after this job is done. We also assume that a task's execution time C_i is always constant.

It's immediately clear that some of these assumptions are not realistic in an actual real-time system. We'll study later on some extensions that can handle deviations of these assumptions.

3.2 Scheduling algorithms

3.2.1 Static priority scheduling algorithms

Static priority scheduling algorithms are algorithms in which each task is assigned a constant priority. The priorities are normally set before executing the a static priority scheduling-algorithm.

3.2.2 Dynamic priority scheduling algorithms

Dynamic priority scheduling algorithms are algorithms in which each task has a priority that may change during the execution of a task set on the CPU.

3.2.3 Preemptive priority scheduling algorithms

A preemptive priority scheduling algorithm is characterized by the fact that these type of algorithms may interrupt the execution of a task to prioritize the execution of another task on the CPU. The task that had been interrupted may resume afterwards.

3.2.4 Rate-monotonic scheduling

A rate-monotonic scheduling algorithm is based on two rules

1. The task with the smallest period has the highest priority,
2. A higher-priority task ready to be executed, overrides the current executed task. The current executed task is interrupted and may resume afterwards.

From the first rule, we can conclude that RMS is a static priority scheduling algorithm because the priorities are assigned to the task before the task set is executed. The second rule implies that RMS is a preemptive priority scheduling algorithm because RMS stops immediately a task with lower priority (higher period) when a task with higher priority (lower period) is ready to be executed, so RMS is a *static-priority scheduling algorithm*.

3.3 Schedulability tests

Given the execution time C_i , the period T_i of a task τ_i , the utilization of the CPU can be calculated as:

$$U_i = \frac{C_i}{T_i}.$$

The utilization factor of a task set \mathcal{T} is defined as:

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \quad (1)$$

When $U > 1$, it's immediately clear that the task set \mathcal{T} cannot be executed by any scheduling algorithm (considering we only have 1 CPU).

3.3.1 Liu & Layland lower bound

Theorem 3.6. *Schedulability test 1 for RMS* Liu & Layland (1973) showed that for n -tasks with distinct periods, a feasible schedule will always exist if the utilization factor of the task set is below $n(2^{\frac{1}{n}} - 1)$:

$$U \leq n(2^{\frac{1}{n}} - 1) \quad (2)$$




Figure 3.1: The lower bound of the schedulability test

Using this theorem, we can easily find that

$$\lim_{n \rightarrow \infty} n(2^{\frac{1}{n}} - 1) = \ln 2 \approx 0.693147... \quad (3)$$

Thus 0.693147... is a worst case bound for RMS. However, this approximation is only a good approximation when the number of tasks is high enough (> 20 tasks). If equation 2 is satisfied, all of the task will meet their deadlines.

Example 3.7. Consider these 3 tasks:

1. τ_1 : $T_1 = 40$ ms, $C_1 = 5$ ms,
2. τ_2 : $T_2 = 40$ ms, $C_2 = 7$ ms,
3. τ_3 : $T_3 = 5$ ms, $C_1 = 2$ ms,

Then: $U_1 = 0.125$, $U_2 = 0.175$, $U_3 = 0.4$. With $n = 3$, the sum of U_1, U_2, U_3 must be lower than 0.7798, using equation 2. We become that $0.125 + 0.175 + 0.4 = 0.7$ is indeed under the limit of 0.7798. By the schedulability test, we conclude that these tasks are schedulable.

3.3.2 Response time (RT) test

Although the schedulability test may be useful, if the utilization function is between the worst case bound and 1, the schedulability test for RMS is inconclusive. A more precise test must be used. This will be *the response time (RT) test*. This test is based on “worst-case phasing”. A worst-case phasing corresponds to a set of values for I_i ($i \in [1..n]$, with a task set of n tasks) leading to the worst schedulable setting of \mathcal{T} . We now have the following nice theorem:

Theorem 3.8. (*Response time (RT) test*) *For a task set, if each task meets its deadline with worst case task phasing, the deadline will always be met.*

When does this worst case task phasing occur? Intuitively, it’s clear that the worst case occurs when all the task τ_i have their phasing equal to 0, so $I_i = 0$. This is the case when all tasks are launched at the same time. This is obvious: take a random task τ_j . In this scenario, the end time $= I_j + T_j = 0 + T_j = T_j$. Now, letting τ_j start at a different moment than the other task, is equal to changing the period of task τ_j to $I'_j > 0$ extends the end time immediately.

Theorem 3.9. *The worst-case phasing occurs when $I_i = 0, \forall i \in [1..n]$. This configuration where all tasks start at the same time is often called a critical instance.*

Now, remember the assumption that a task can only have one job that’s being executed, new jobs have to wait after this job is done. This immediately leads to the following theorem, a reformulation of the RT test:

Theorem 3.10. (*Reformulation - Response time (RT) test*) *A task set can be scheduled by RMS if the deadline of the first job of each task is met when using the scheduling algorithm starting from a critical instant.*

Remark 3.11. *Theorem 3.10 applies in a much more general setting too: a periodic task set can be scheduled in any static scheduling algorithm if the deadline of each task is met when using the scheduling algorithm starting from a critical instant.*

3.4 Time-Demand function

Based on the above test, we conclude that the RT test requires the computation of the total processing time of each task in a task set. If each

total processing time is less than it's corresponding period, the task set is schedulable.

The total processing requirement $u_i(t)$ of a task τ_i in the time interval $[0, t]$ is given by, with $0 < t \leq T_i$ (note that task are ordered by increasing priorities!):

$$u_i(t) = C_i + \sum_{k=1}^{i-1} \left\lceil \frac{t}{T_i} \right\rceil C_k \quad (4)$$

The idea is immediately clear: if $u_i(t) \leq t$ for some $t \leq T_i$ then task τ_i is schedulable.

Example 3.12. Consider these 3 tasks (for the shake of simplicity, we used very, very short tasks):

1. τ_1 : $T_1 = 4$ ms, $C_1 = 1$ ms,
2. τ_2 : $T_2 = 5$ ms, $C_2 = 2$ ms,
3. τ_3 : $T_3 = 7$ ms, $C_3 = 2$ ms,

Remember from 3.3.1 that we first have to check the schedulability test with the lower bound, then: $U_1 = 0.25$, $U_2 = 0.4$, $U_3 = 0.28$. With $n = 3$, the sum of U_1, U_2, U_3 must be lower than 0.7798, using equation 2. We become that $0.25 + 0.4 + 0.28 = 0.91 > 0.7798$. So we conclude that can not tell us if this task set is schedulable as a whole.

We thus preform a time demand analysis, then:

1. $u_1(t) = C_1 = 1$
2. $u_2(t) = C_2 + \left\lceil \frac{t}{T_1} \right\rceil C_1 = 2 + \left\lceil \frac{t}{4} \right\rceil * 1$
3. $u_3(t) = C_3 + \left\lceil \frac{t}{T_1} \right\rceil C_1 + \left\lceil \frac{t}{T_2} \right\rceil C_2 = 2 + \left\lceil \frac{t}{4} \right\rceil * 1 + \left\lceil \frac{t}{5} \right\rceil * 2$

We test:

1. $u_1(t) \leq t$ satisfied for $t = 4$? $\rightsquigarrow u_1(4) = 1 \leq 4 \Rightarrow \mathbf{OK!}$
2. $u_2(t) \leq t$ satisfied for $t \in 4, 5$? $\rightsquigarrow u_2(4) = 3 \leq 4, u_2(5) = 4 \leq 5 \Rightarrow \mathbf{OK!}$
3. $u_3(t) \leq t$ satisfied for $t \in 4, 5, 7$? $\rightsquigarrow u_3(4) = 5 > 4, u_3(5) = 6 > 5, u_3(7) = 8 > 7 \Rightarrow \mathbf{NOT OK!}$

We conclude that τ_1 and τ_2 are schedulable, but τ_3 isn't.

So, to perform a time demand analysis in general: we have to do the following steps:

Method 3.13. Time Demand Analysis

1. For each $i \in [1..n]$, calculate the time demand function $u_i(t)$:

$$C_i + \sum_{k=1}^{i-1} \left\lceil \frac{t}{T_i} \right\rceil C_k$$

2. Check whether the inequality $u_i(t) \leq t$ is satisfied for values of t that are equal to

$$t = j * p_k$$

$$\text{with } k = 1, 2, \dots, i \text{ and } j = 1, 2, \dots, \left\lfloor \frac{T_i}{T_j} \right\rfloor$$

The time complexity of the time-demand analysis for each task is $O(n(\frac{T_n}{T_1}))$

In most references, they will use the notion of the cumulative demand on a CPU of a task set \mathcal{T} until time t , this is given by:

Definition 3.14. Cumulative demand The cumulative demand on a CPU of a task set $\mathcal{T} = \{\tau_1, \tau_2, \dots, \tau_n\}$ is defined as the computation time that is needed to achieve the current jobs of task τ_1, \dots, τ_i and the previous jobs of the same task.

$$U_i(t) = \sum_{k=1}^i u_k(t) = \sum_{k=1}^i C_k \left\lceil \frac{t}{T_k} \right\rceil$$

3.5 Dynamic scheduling variant of RMS

Liu and Layland also studied an *earliest deadline first*-algorithm: instead of executing the task with the highest priority first, this algorithm will choose the task with the earliest deadline first. This variant of RMS is a dynamic scheduling algorithm because deadlines can change during the execution of a task set on a CPU. The calculation of the earliest deadline is calculated every time a job is finished.

When the utilization factor of a task set \mathcal{T} equals 1, this algorithm will use 100% of the capacity of the CPU. This seems really powerful, as we know from 3.3.1 that the rate monotonic scheduling has a (worst case) bound of 69%. One may wonder why we don't pay much more attention to this algorithm.

This has plenty of good reasons: first, a CPU usage of 100% is only achieved when a task set is easy to schedule. Also there are some serious stability issues with this algorithm: a system with a CPU usage of 100% can

easily be overloaded. This can lead to essential tasks missing their deadlines. Using RMS, we know for sure this will work out as the essential tasks have the highest priority and are executed first. This is not possible with the *earliest deadline first*-algorithm, priority of tasks is not taken into consideration. Also, dynamic scheduling algorithms have some overload: they need runtime support to calculate which task must be executed next. Moreover, the difference is in practice rather small. RMT can achieve a CPU-usage as high as 90%.

4 Extensions

In 3.1 we made some assumptions about the ‘nature’ of our tasks. With these assumptions we get a configuration in which it is easy to reason about rate monotonic scheduling. But as research evolves, the rate monotonic theory was extended for other configurations with other assumptions too. Today it is possible to use rate monotonic theory for both periodic and non-periodic tasks with synchronization requirements. We now take a look at some of these extensions of rate monotonic scheduling.

4.1 Non-periodic services

Most systems have to deal with non-periodic tasks. Non-periodic tasks are tasks with dynamic intervals between requests. These tasks can be included in the rate monotonic theory by adding of one or more non-periodic services. A non-periodic service will handle arriving requests based on their assigned priority as long as there is execution budget available. The assigned priority is decided by the rate monotonic algorithm based on the completion period.

4.1.1 Sporadic service

A sporadic service is a conceptual task with an *execution budget* and a *completion period*. A sporadic service preserves its execution time until a non-periodic job occurs. The completion time is not replenished periodically, but is replenished only after a period in which all of the execution time is used by a non-periodic task. The completion time is set at first on the current time added by the period of the sporadic service. The execution time that is replenished equals the execution time consumed since the last execution of the sporadic service.

This implementation avoids the deferred execution effect, that was a big problem in earlier solutions. With this problem avoided, we immediately see that the sporadic service is a non-periodic service behaving like any periodic task under the rate monotonic assumptions. This is not the case for the most other solutions preserving the execution time.

However, some caution is appropriate. A sporadic service can be viewed as a periodic task so it’s very attractive to use the known tests to decide about the schedulability of a task set containing some non-periodic tasks. But we have to keep in mind that one or more sporadic services must be created to handle this task set. It is possible to guarantee that the non-periodic tasks will meet their deadlines but a small interval between two requests is needed.

If the period of the sporadic service is less than the deadline of a non-periodic task, this method can't be used.

4.2 Priority ceiling protocol

4.2.1 Priority inversion

Priority inversion occurs when a high priority task with a hard deadline may be blocked by a low priority task that uses a shared resource. We have two types of priority inversion: bounded and unbounded. They both occur when two tasks want to use a shared resource. The period of time a task is locked to use a shared resource is called a tasks critical region.

4.2.2 The protocol

When we want to apply rate monotonic analysis to systems where tasks share resources, we have to find a solution for *priority inversion*. This can cause mutual locks between tasks, eventually leading to a deadlock. Now, two key rules are added to avoid this:

Definition 4.1. (*Priority ceiling*) *A semaphore gets a priority ceiling equal to the highest priority of all tasks that may use this semaphore.*

Definition 4.2. (*Priority inheritance*) *A task τ_1 is only allowed to enter a semaphore S if the semaphore is not already in use and if the task's priority is higher than its inherited priority ceilings of all the other semaphores that are locked at the same moment. Otherwise the task is temporary blocked.*

By adding this two rules, mutual locks and deadlocks are avoided. With these guarantees, RM analysis allows to analyze tasks where sources are shared. To incorporate this into the mathematical reasoning from the previous section, terms of blocking must be added to the reasoning. Let B_i the maximum blocking time for task τ_i , we can, for example, reformulate 3.3.1 into:

Theorem 4.3. *Schedulability test for RMS, Lower Bound - adjusted for the priority ceiling protocol* *A feasible schedule will always exist with respect to the priority ceiling protocol if:*

$$U \leq n(2^{\frac{1}{n}} - 1) + \sum_{j=1}^n \frac{B_j}{T_j} \quad (5)$$

for the task set $\mathcal{T} = \{\tau_1, \dots, \tau_n\}$

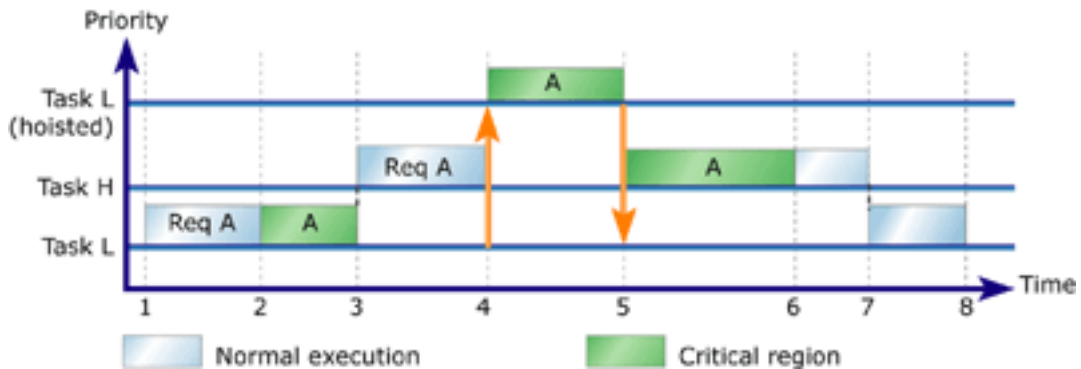


Figure 4.2: An example of the priority ceiling protocol

Example 4.4.¹

Look at figure 4.2. The explanation:

1. Task L receives control of the processor and begins executing.
 - (a) The task makes a request for Resource A.
2. Task L is granted ownership of Resource A and enters its critical region.
3. Task L is preempted by Task H, a higher-priority task.
 - (a) Task H begins executing and requests ownership of Resource A, which is owned by Task L.
4. Task L is hoisted to a priority above Task H and resumes executing its critical region.
5. Task L releases Resource A and is lowered back to its original priority.
 - (a) Task H acquires ownership of Resource A and begins executing its critical region.
6. Task H releases Resource A and continues executing normally.
7. Task H finishes executing and Task L continues executing normally.
8. Task L finishes executing.

¹This example and the related picture is retrieved from <http://www.embedded.com/design/configurable-systems/4024970/How-to-use-priority-inheritance>

4.3 Tasks with $D_i \leq T_i$

We also made the assumption that the deadline D_i of a task τ_i is the same as his period T_i . Dropping this assumption will lead to more complex formulas. The Rate-Monotonic scheduling algorithm doesn't depend on this, only the schedulability tests must be reformulated. We reformulate the schedulability test 1 for RMS as an example.

Theorem 4.5. *Schedulability test 1 for RMS - adjusted for $D_i \leq T_i$*
A feasible schedule will always exist for a task set nif

$$\forall j \in [1...n] : \sum_{i=1}^j \frac{C_i}{T_i} \leq \mathcal{U} \left(\frac{D_j}{T_j} \right) \quad (6)$$

with

$$\mathcal{U} \left(\frac{D_j}{T_j} \right) = \begin{cases} 2j^{\frac{D_j}{T_j} - \frac{1}{j}} - j + \frac{D_j}{T_j} + 1 & \text{if } \frac{1}{2} \leq \frac{D_j}{T_j} \leq 1 \\ \frac{D_j}{T_j} & \text{if } 0 \leq \frac{D_j}{T_j} \leq \frac{1}{2} \end{cases} \quad (7)$$

If $D_j = T_j$, then $\frac{D_j}{T_j} = 1$, we become that $\mathcal{U}(1) = j(2^{\frac{1}{j}} - 1)$, this exactly the 'normal' formula we became in 3.3.1!

The RT-test is adjusted in the same way.

4.4 Other extensions

There are a lot more extensions for rate monotonic analysis to make it more widely applicable. By example [3] explains how to deal with context switching overhead an preemption by fixed priority interrupt task. [4] explains a protocol for a rate monotonic scheduling algorithm that can handle changes (addition, deletion) of tasks during the execution of a task set. There is also research about statistical rate-monotonic scheduling.

5 Simulations using the schedulability analysis tool Cheddar

5.1 What is Cheddar?

There are a number of tools available where a developer can do a schedulability analysis and a calculation of response times. Initially, I planned to use rapidRMA from Tri-Pac software, but you need a free license key for this software package that you can obtain by sending them an e-mail. Unfortunately, after two weeks I still had no answer. I decided to use Cheddar instead. Cheddar is a free real time scheduling tool designed for checking task temporal constraints of a real time application/system. Cheddar is an open-source tool developed by members from the LISyC laboratory of the Université de Bretagne Occidentale and Ellidiss Technologies. It can be downloaded from <http://beru.univ-brest.fr/~singhoff/cheddar/>.

5.2 Configuration

To configure the simulator Cheddar, I refer to the install package itself: in the folder `/project_examples/xml` you can find very clear examples that are straightforward to change. The idea is that you first define a configuration (which processors, memory, tasks, resources,...).

5.3 Simulation of RMS of example 5.3

We retake example and will simulate it by Cheddar. We define a single processor and the RMS scheduler. In the next figures, some screenshots are given. The result you can find in 5.4

1. τ_1 : $T_1 = D_1 = 4$ ms, $C_1 = 1$ ms,
2. τ_2 : $T_2 = D_2 = 5$ ms, $C_2 = 2$ ms,
3. τ_3 : $T_3 = D_3 = 7$ ms, $C_3 = 2$ ms,

Just as in our calculations in the example, we become indeed that task 3 will miss its deadline. The task set is not schedulable.

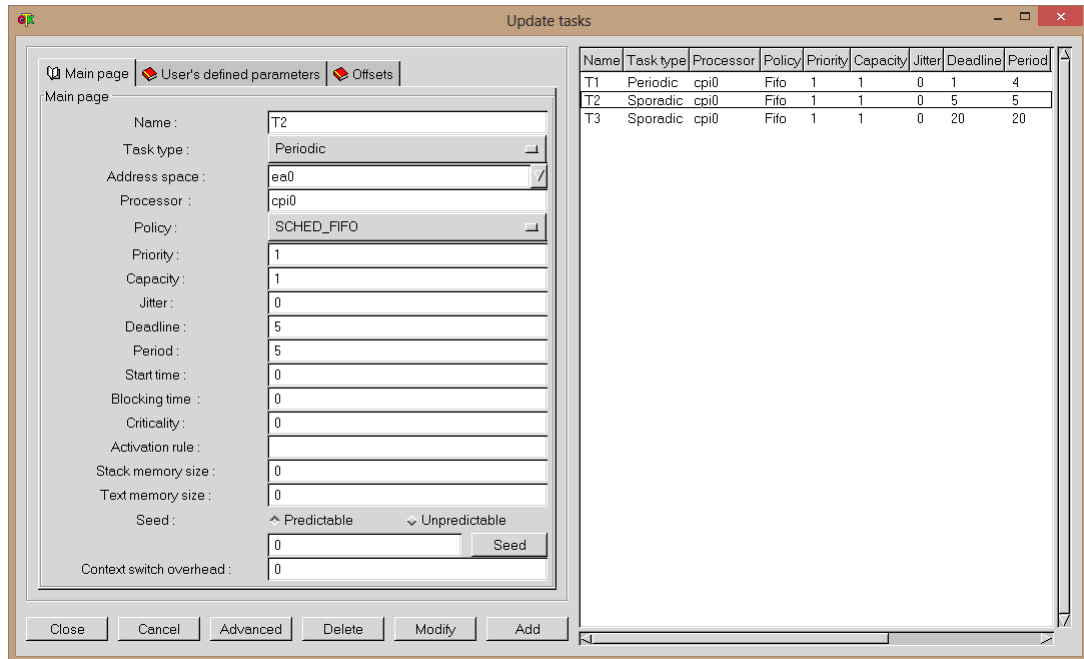


Figure 5.3: Adding the tasks for the simulation.

6 Conclusion

This paper gave an overview of rate-monotonic scheduling. In the first section, we contextualize the subject in the area where RMS is used. Next, the algorithm and all the needed terminology and definitions were introduced in an easy configuration with a lot of assumptions on the tasks. We also introduced the schedulability tests needed to analyze whether a certain task set can be scheduled with rate-monotonic scheduling or not. In the next section, we dived into some extensions. In the first chapter we made a lot of assumptions about the tasks being executed by RMS, but now we take a look at RMS-adaptations applied on other settings, often these are much more general cases. The last chapter finally gives some concrete examples of simulations of RMS.

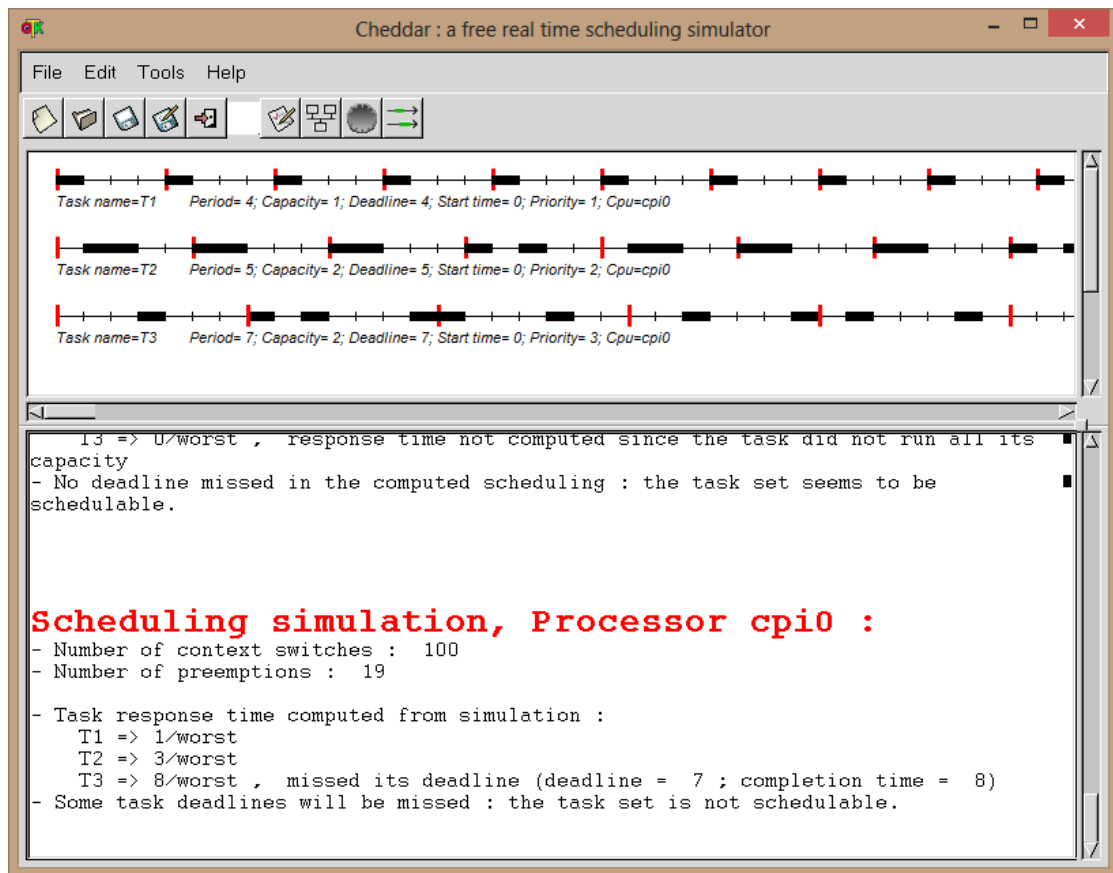


Figure 5.4: The simulation with Cheddar gives the expected result.

References

- [1] Fowler, Priscilla., & Levine, Linda. (1993). *Technology Transition Push: A Case Study of Rate Monotonic Analysis (Part 1)* (CMU/SEI-93-TR-029). Retrieved January 06, 2014, from the Software Engineering Institute, Carnegie Mellon University.
- [2] Jones, Michael. How to user priority inheritance, retrieved from, <http://www.embedded.com/design/configurable-systems/4024970/How-to-use-priority-inheritance>, on 7, january 2013
- [3] Obenza, Ray, and Mendal, Geoff. *Guaranteeing Real Time Performance Using RMA*, The Embedded Systems Conference, San Jose, CA, 1998
- [4] Sha, L., R. Rajkumar, and S. Sathaye. Priority inheritance protocols: An approach to real-time synchronization. *IEEE Transactions on Computers*, 39 (9), 1175-1185: 1990.
- [5] Sha, Lui, Klein, Mark H., and Goodenough, John B. *Rate Monotonic Analysis*, Technical Report CMU/SEI-91-TR-6 ESD-91-TR-6, March 1991.
- [6] F. Singhoff, J. Legrand, L. Nana, L. Marc'Al. *Cheddar : a Flexible Real Time Scheduling Framework*. ACM SIG Ada Ada Letters, volume 24, number 4, pages 1-8. Edited by ACM Press, New York, USA. December 2004, ISSN:1094-3641.
- [7] F. Singhoff, J. Legrand, L. Nana, L. Marc'Al. *Extending Rate Monotonic Analysis when Tasks Schare Buffers*. University of Brest.
- [8] ,Nate Forman. *Rate Monotonic Theory* .
- [9] Frank Drews,. *Rate Monotonic Scheduling* .